

---

# Analysis on Logistic and Softmax Regression Using MNIST Dataset

---

<b>Chetan Gandotra</b>	<b>Rishabh Misra</b>
UC San Diego	UC San Diego
9500 Gilman Drive	9500 Gilman Drive
cgandotr@ucsd.edu	r1misra@ucsd.edu

## Abstract

1 This report discusses the first programming assignment of course CSE 253: Neural  
2 Networks and Pattern Recognition, its solutions and the inferences. MNIST dataset  
3 was used and the hand-written digits in it were classified using Logistic and Softmax  
4 Regressions. Under Logistic Regression, two-way classification was performed  
5 on specific digits (2's and 3's, 2's and 8's). An accuracy of more than 97% was  
6 achieved on both of these subsets of data using Logistic Regression. For Softmax  
7 Regression, we performed a ten-way classification (for all digits from 0 to 9) and  
8 achieved an accuracy of 87.65% on the test set.

## 9 1 Derivation of Gradient for Logistic Regression

### 10 1.1 Introduction

The problem statement here is to find the gradient of the cost function. The error function for the logistic regression follows from the negative log likelihood, which can be written as:

$$E(w) = - \sum_{n=1}^N \{t^n \ln y^n + (1 - t^n) \ln(1 - y^n)\}$$

### 11 1.2 Methodology

In this section, we will derive the gradient of cost function, which will be used in the later parts of this report. To find the optimal weight parameters, we need to take the partial derivative of the error function with respect to  $w_j$  as follows:

$$\frac{\partial E(w)}{\partial w_j} = - \sum_{n=1}^N \left[ t^n \frac{\partial \ln y^n}{\partial w_j} + (1 - t^n) \frac{\partial \ln(1 - y^n)}{\partial w_j} \right]$$

$$\frac{\partial E(w)}{\partial w_j} = - \sum_{n=1}^N \left[ \frac{t^n}{y^n} \frac{\partial y^n}{\partial w_j} + \frac{(1 - t^n)}{1 - y^n} \frac{\partial (1 - y^n)}{\partial w_j} \right]$$

Since  $y^n = \sigma(\mathbf{w} \cdot \mathbf{x}^n)$ , the derivative can be written as:

$$\frac{\partial E(w)}{\partial w_j} = - \sum_{n=1}^N \left[ \frac{t^n}{\sigma(\mathbf{w} \cdot \mathbf{x}^n)} \frac{\partial \sigma(\mathbf{w} \cdot \mathbf{x}^n)}{\partial w_j} + \frac{(1 - t^n)}{1 - \sigma(\mathbf{w} \cdot \mathbf{x}^n)} \frac{\partial (1 - \sigma(\mathbf{w} \cdot \mathbf{x}^n))}{\partial w_j} \right]$$

12 Using the following properties of sigmoid function

$$\sigma(-\mathbf{x}) = 1 - \sigma(\mathbf{x}) \quad (1)$$

$$\frac{\partial \sigma(-\mathbf{x})}{\partial x} = \sigma(\mathbf{x})\sigma(-\mathbf{x}) \quad (2)$$

the derivative can be written as:

$$\frac{\partial E(w)}{\partial w_j} = - \sum_{n=1}^N \left[ \frac{t^n}{\sigma(\mathbf{w} \cdot \mathbf{x}^n)} \sigma(\mathbf{w} \cdot \mathbf{x}^n) \sigma(-\mathbf{w} \cdot \mathbf{x}^n) x_j^n - \frac{(1-t^n)}{\sigma(-\mathbf{w} \cdot \mathbf{x}^n)} \sigma(-\mathbf{w} \cdot \mathbf{x}^n) \sigma(\mathbf{w} \cdot \mathbf{x}^n) x_j^n \right]$$

$$\frac{\partial E(w)}{\partial w_j} = - \sum_{n=1}^N x_j^n [t^n \sigma(-\mathbf{w} \cdot \mathbf{x}^n) - (1-t^n) \sigma(\mathbf{w} \cdot \mathbf{x}^n)]$$

Using (1) we get,

$$\frac{\partial E(w)}{\partial w_j} = - \sum_{n=1}^N x_j^n [t^n (1 - \sigma(\mathbf{w} \cdot \mathbf{x}^n)) - (1-t^n) \sigma(\mathbf{w} \cdot \mathbf{x}^n)]$$

Solving the above equation, we get:

$$\frac{\partial E(w)}{\partial w_j} = - \sum_{n=1}^N x_j^n [t^n - \sigma(\mathbf{w} \cdot \mathbf{x}^n)]$$

or

$$-\frac{\partial E(w)}{\partial w_j} = \sum_{n=1}^N (t^n - y^n) x_j^n$$

### 13 1.3 Results

Hence, from the above derivation, it follows that for  $n^{th}$  sample, the gradient of error can be written as:

$$-\frac{\partial E^n(w)}{\partial w_j} = (t^n - y^n) x_j^n$$

### 14 1.4 Discussion

15 The expression takes the difference between true label and predicted label and weigh it by the input  
 16 data value. It makes sense because if there is a stark difference between the true and the predicted  
 17 labels, the gradient value would be large. Thus, the corresponding component of the weight vector  
 18 would be adjusted quickly in the direction of gradient to reduce the loss.

## 19 2 Derivation of Gradient for Softmax Regression

### 20 2.1 Introduction

In this section, the focus is to find the gradient of the loss function of Softmax Regression -  $E(w)$ . The error function for the softmax regression follows from the negative log likelihood, which can be written as:

$$E(w) = - \sum_{n=1}^N \sum_{k'=1}^C t_{k'}^n \ln y_{k'}^n$$

### 21 2.2 Methodology

To find the optimal weight parameters for each class, we need to take the partial derivative of the error function with respect to  $w_{jk}$  as follows:

$$\begin{aligned} \frac{\partial E(w)}{\partial w_{jk}} &= - \sum_{n=1}^N \sum_{k'=1}^C \left[ t_{k'}^n \frac{\partial \ln y_{k'}^n}{\partial w_{jk}} \right] \\ - \frac{\partial E(w)}{\partial w_{jk}} &= \sum_{n=1}^N \left[ \frac{t_k^n}{y_k^n} \frac{\partial y_k^n}{\partial w_{jk}} \right] + \sum_{k' \neq k} \left[ \frac{t_{k'}^n}{y_{k'}^n} \frac{\partial y_{k'}^n}{\partial w_{jk}} \right] \end{aligned} \quad (3)$$

Since,

$$y_k^n = \frac{e^{\mathbf{w}_k \cdot \mathbf{x}^n}}{\sum_{l=1}^C e^{\mathbf{w}_l \cdot \mathbf{x}^n}}$$

22 the derivatives  $\frac{\partial y_k^n}{\partial w_{jk}}$  and  $\frac{\partial y_{k'}^n}{\partial w_{jk}}$  can be written as:

$$\begin{aligned} \frac{\partial y_k^n}{\partial w_{jk}} &= \frac{e^{\mathbf{w}_k \cdot \mathbf{x}^n}}{\sum_{l=1}^C e^{\mathbf{w}_l \cdot \mathbf{x}^n}} x_j^n - e^{\mathbf{w}_k \cdot \mathbf{x}^n} \left[ \frac{1}{(\sum_{l=1}^C e^{\mathbf{w}_l \cdot \mathbf{x}^n})^2} \right] e^{\mathbf{w}_k \cdot \mathbf{x}^n} x_j^n \\ \frac{\partial y_k^n}{\partial w_{jk}} &= (y_k^n - (y_k^n)^2) x_j^n \end{aligned} \quad (4)$$

$$\begin{aligned} \frac{\partial y_{k'}^n}{\partial w_{jk}} &= -e^{\mathbf{w}_{k'} \cdot \mathbf{x}^n} \left[ \frac{1}{(\sum_{l=1}^C e^{\mathbf{w}_l \cdot \mathbf{x}^n})^2} \right] e^{\mathbf{w}_{k'} \cdot \mathbf{x}^n} x_j^n \\ \frac{\partial y_{k'}^n}{\partial w_{jk}} &= -(y_{k'}^n)^2 x_j^n \end{aligned} \quad (5)$$

23 Substituting (4) and (5) in (3), we get

$$\begin{aligned} - \frac{\partial E(w)}{\partial w_{jk}} &= \sum_{n=1}^N \left[ \frac{t_k^n}{y_k^n} (y_k^n - (y_k^n)^2) x_j^n \right] - \sum_{k' \neq k} \left[ \frac{t_{k'}^n}{y_{k'}^n} (y_{k'}^n)^2 x_j^n \right] \\ - \frac{\partial E(w)}{\partial w_{jk}} &= \sum_{n=1}^N [t_k^n (1 - y_k^n) x_j^n] - \sum_{k' \neq k} [t_{k'}^n y_{k'}^n x_j^n] \end{aligned} \quad (6)$$

24 Now, for any sample, only one of the  $C$  labels in  $t^n$  would be 1, and all the others would be 0. This is  
25 because the label one would be the label set, and each training example can correspond to only one  
26 label. Thus, for any sample  $a$  where  $t_k^n$  is 1, the derivative would be:

$$-\frac{\partial E^a(w)}{\partial w_{jk}} = [(1 - y_k^a) x_j^a]$$

27 or it could be written as:

$$-\frac{\partial E^a(w)}{\partial w_{jk}} = [(t_k^n - y_k^a) x_j^a] \quad (7)$$

For any sample  $b$  where one of  $t_k^n$  is 1 (where  $k' \neq k$ ), the derivative would be:

$$-\frac{\partial E^b(w)}{\partial w_{jk}} = [(-y_k^b) x_j^b]$$

28 or it could be written as:

$$-\frac{\partial E^b(w)}{\partial w_{jk}} = [(t_k^n - y_k^b) x_j^b] \quad (8)$$

29 Using the results of (7) and (8), (6) could be written as:

$$-\frac{\partial E(w)}{\partial w_{jk}} = \sum_{n=1}^N (t_k^n - y_k^n) x_j^n$$

### 30 2.3 Results

31 Thus, using our findings above, we can say that for  $n^{th}$  sample, the derivative can be written as:

$$-\frac{\partial E^n(w)}{\partial w_{jk}} = (t_k^n - y_k^n) x_j^n \quad (9)$$

### 32 2.4 Discussion

33 Interestingly, the expression of gradient looks similar to that of logistic regression. In this case, the  
 34 derivative takes the difference between true label and predicted label for the  $k^{th}$  class and weigh it  
 35 by the input data value. Again, if the difference is big between the true and the predicted labels, the  
 36 gradient value would be large. Thus, the corresponding component of the weight vector would be  
 37 adjusted quickly in the direction of gradient to reduce the loss.

### 3 Read in Data

#### 3.1 Introduction

As mentioned in the abstract, we deal with the "MNIST" dataset in this programming assignment. The first and foremost task before operating on the data was to load it.

#### 3.2 Methodology

The MNIST data was downloaded from the website at <http://yann.lecun.com/exdb/mnist/> (the same link as given in PA1). To read the data, GitHub library at <https://github.com/akosiorek/CSE/blob/master/MLCV/> was used, which returns the training and testing data in matrix form, and labels as a vectors. Operations were then performed on this data to add a column of ones (bias term) at the beginning and to extract digit specific data (2-3's and 2-8's). Also, the data was restricted to first 20k entries, 10% of which was allocated to a hold-out set. The size of test data was kept as two thousand. This was done by picking the first 2k entries from the test data returned by the library.

#### 3.3 Results

Using some existing libraries on Github, we were able to extract the data into variables in Python. This data, however, consisted of the full 60k training data points and 10k testing data points. We extracted the first 20k training data points and the first 2k testing data points. 10% of the training data was designated as a hold-out set.

#### 3.4 Discussion

Computation on large data sets can often be time consuming. Due to this reason, we extracted the full data and restricted the size of training, validation and test sets. This allowed faster computations throughout the programming assignment. A hold-out set acts as a dummy test set, which we use so as to improve performance of our model by testing it on the hold-out set. Good accuracy on hold-out set leads to a good accuracy on test set in general.

## 4 Logistic Regression via gradient descent

#### 4.1 Introduction

In this part, we are required to use logistic regression and classify a given hand-written digit as either 2 or 3. Since logistic regression uses binary output, we say that the target is 1 if the input is from the "2" category and 0 if it is from the other category. We are required to produce the following:

1. Plot of loss function ( $E$ ) over the training set, test set and the hold-out set
2. Plot of percent correct classification over training for the training set, the hold-out set, and the test set
3. The above two plots for digits 2 and 8
4. Display weights as images for both the classifiers (2 vs. 3 and 2 vs. 8). Plot the difference between weights as well.

#### 4.2 Methodology

To plot the first graph, loss function is put against the Y-axis, and the iteration number along the X-axis. The loss function used was:

$$E(w) = - \sum_{n=1}^N \{t^n \ln y^n + (1 - t^n) \ln(1 - y^n)\}$$

This was done for all the three sets - training, test and validation. Hence, for each set, the value of  $N$  and the corresponding data/labels change. This procedure was repeated for the 2's and 8's data set.

In the next graph, we plotted "percent correct" against the iteration number. The value of percent correct can be inferred by going over all the examples as test set in a way, and seeing what our model predicts on it. For every correct classification, we add one to the number of data points classified correctly. Then, at the end, we find the corresponding percentage. This is repeated at each iteration for all the three sets - training, test and validation. Again, this procedure was repeated for the 2's and 8's data set.

To display the weight vectors for both the cases and their difference, the bias term in them was dropped. This reduced the weight vector to a 784 dimensional vector, which was re-shaped into a 28 x 28 matrix and then plotted using Python.

### 4.3 Results

Various results are plotted below:

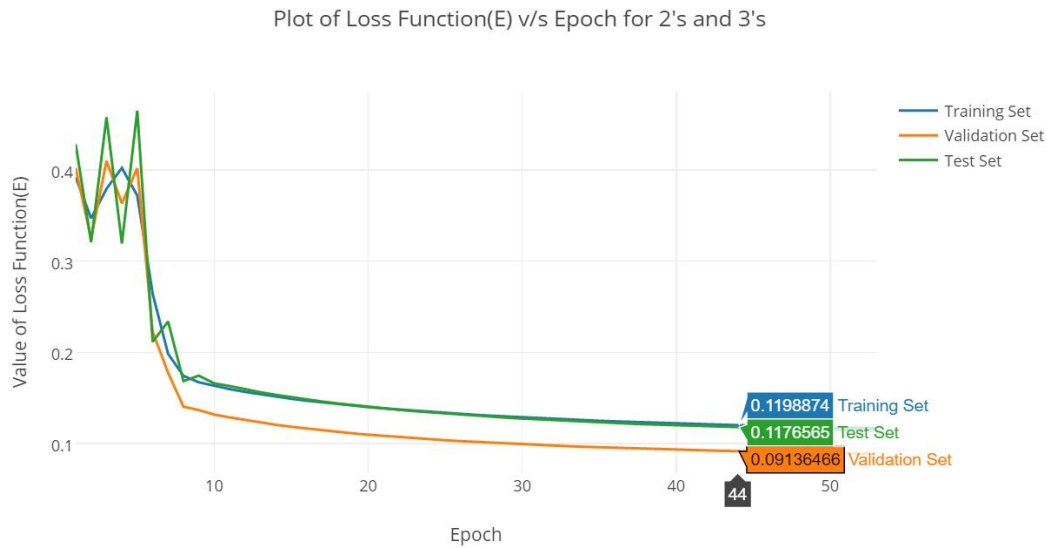


Figure 1: Loss Function (E) vs Epoch for 2's and 3's



Figure 2: Loss Function calculated every 1/20 of Epoch vs Epoch for 2's and 3's

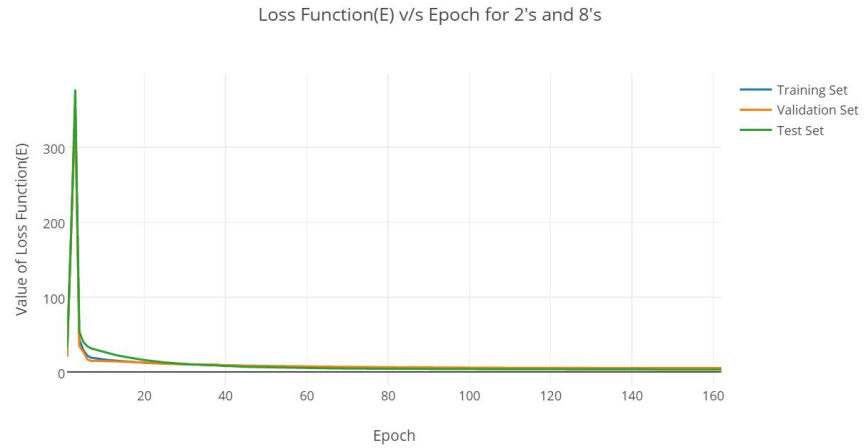


Figure 3: Loss Function (E) vs Epoch for 2's and 8's

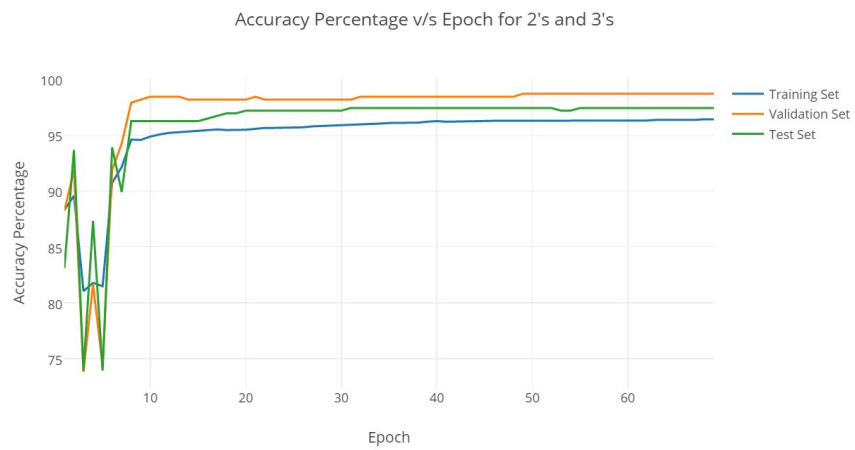


Figure 4: Percent correct classification vs Epoch for 2's and 3's

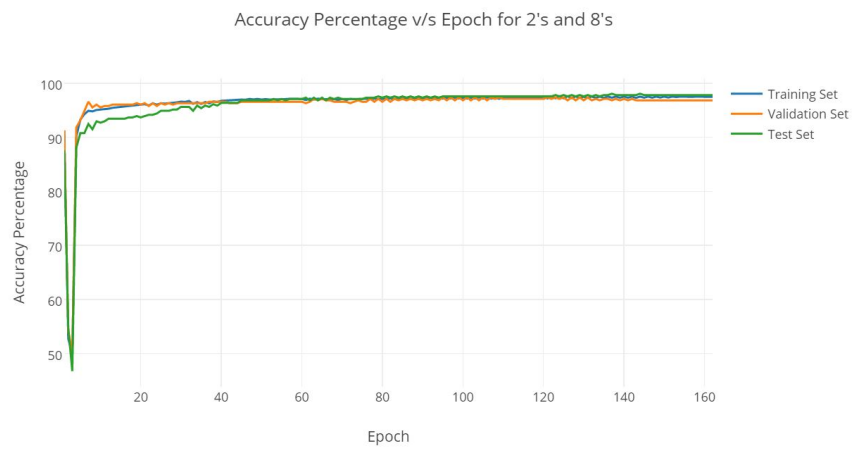


Figure 5: Percent correct classification vs Epoch for 2's and 8's




Figure 6: Weights as an image for 2's and 3's




Figure 7: Weights as an image for 2's and 8's




Figure 8: Difference of weights as an image for 2's and 3's and 2's and 8's

#### 87 4.4 Discussion

88 Generally, we assume that the data for training and testing are generated from same underlying  
 89 distribution. Thus, they have similar underlying properties. However, since test data is not available  
 90 to us in the real world, we extract a portion of training data to test our model and call it a validation  
 91 set. And since the validation set and the test set are generated from the same underlying distribution,  
 92 the performance of our model should not vary much on validation and test set in ideal case (it would  
 93 vary a little bit in our case as validation set is small as compared to test set). The aforementioned  
 94 behavior is reaffirmed by our experiments. From figure ?? and figure ??, we see that curves for  
 95 validation error and test error are similar and validation error hovers around the test error. Same is the  
 96 case for classification of 2's and 8's in figure ??. Thus, we can safely say that validation set closely  
 97 captures how our model would perform on test data. One interesting thing to note is that in the case  
 98 of batching (contained 1/20 of all points), the curve of error function is smooth . This is because  
 99 the gradient is calculated only on few points at a time, which makes the gradient increase gradually  
 100 towards the optimum.

101 We have chosen the value of the hyperparameter T as 2000 empirically. Also, for the early stopping  
 102 of our algorithm, we have checked whether the error is non-decreasing on 15 iterations. The reason  
 103 for choosing a bit large value is that the error on the validation set used to change in the steps of 5-10  
 104 iterations. Thus, just to insure that the error does not decrease after early stopping, we took added  
 105 some buffer iterations.

106 In figure ?? and figure ??, we see that the accuracy percentage increases over time. This is expected,  
 107 as with more number of iterations, we minimize loss and fit the data. This means that we will be able



108 to classify more number of points correctly leading to higher accuracy. Also, note that the graph of  
109 2's and 8's is more overlapping. This can be because the validation/test data is very similar to the  
110 training data.

111 Next, we plot the weight vectors for three cases - 2's and 3's, 2's and 8's, and the difference of these  
112 two in figure ??, ?? and ?? respectively. For both 2's and 3's case and the 2's and 8's case, the weight  
113 vector seems like images of 2 and 3/8 have been superimposed. This seems intuitively correct as  
114 the weight vector needs to predict either of these. For the difference case, the image looks like a  
115 mirror-image of 3. This is because the 2's component is common in both the cases and must get  
116 cancelled out. Thereafter, we are left with 8's and 3's, in which superimposition cancels out too.  
117 Hence, what we are left with is the portion of 8 that was not covered by 3 and thus looks like a mirror  
118 image of 3.

## 119 5 Regularization

### 120 5.1 Introduction

121 Regularization is a commonly used technique to improve the model generalization. We write the  
122 regularized loss function  $J(w)$  as:

$$J(w) = E(w) + \lambda C(w)$$

123 where  $C(w)$  is the complexity penalty and  $\lambda$  is the strength of regularization. For  $\lambda = 0$ ,  $J$  reduces to  
124  $E$ . Considering  $L_2$  norm as the complexity penalty, we have:

$$J(w) = E(w) + \lambda ||w||^2$$

125 For  $L_1$  norm, we have:

$$J(w) = E(w) + \lambda |w|$$

126 In the first part, we are expected to derive the update term for both  $L_1$  and  $L_2$  penalties.

127 Next, we are expected to plot the percent correct v/s iterations graph for different  $\lambda$  values. This is  
128 followed by plotting length of weight vector v/s iterations for different  $\lambda$  values. Then, we plot the  
129 final test error with each of the  $\lambda$ . Finally, we are expected to plot the weights as images.

### 130 5.2 Methodology

131 To derive the update term, we take derivative of this function with respect to  $w$ , the weight vector.  
132 Hence, we have:

$$\frac{\partial J(w)}{\partial w} = \frac{\partial E(w)}{\partial w} + \lambda \frac{\partial C(w)}{\partial w}$$

133 We have already calculate the first part of the equation -  $\frac{\partial E(w)}{\partial w}$  in Question 1. Hence, according to  
134 the question, solving for  $\frac{\partial C(w)}{\partial w}$ , we have:

$$\frac{\partial C(w)}{\partial w} = \frac{||w||^2}{\partial w} = 2w$$

Therefore, we have:

$$\frac{\partial J(w)}{\partial w} = \frac{\partial E(w)}{\partial w} + 2\lambda w$$

Similarly, for  $L_1$  norm as the complexity penalty, we have

$$J(w) = E(w) + \lambda |w|$$

Therefore,

$$\frac{\partial J(w)}{\partial w} = \frac{\partial E(w)}{\partial w} + \frac{\partial \lambda |w|}{\partial w}$$

$$\frac{\partial J(w)}{\partial w} = \frac{\partial E(w)}{\partial w} + \lambda \frac{\partial |w|}{\partial w}$$

135 Now, the entry at index  $j$  of partial derivative of  $|w|$  can be written as:

$$\frac{\partial |w|}{\partial w_j} = \begin{cases} 1, & \text{if } w_j \geq 0 \\ -1, & \text{otherwise} \end{cases}$$

Thus, the value of  $\frac{\partial |w|}{\partial w}$  is A vector of all one's or minues one's depending upon the sign of entries in (w) vector and has the same number of elements as in w.

To plot the graphs, a similar approach as in Section 4 was undertaken. The only difference was that earlier we did it for different data sets - training, test and validation. In these graphs, we always take the training set and calculate percent error and length of weight vector at that iteration. This is done multiple times by changing  $\lambda$  values. Then, we plot the final test error for each  $\lambda$  value, keeping the learning rate fixed. This plot is made as a bar graph, with one bar for each  $\lambda$ . Finally, we plot the weights as images like we did in Section 4. To display the weight vectors, the bias term was dropped. This reduced the weight vector to a 784 dimensional vector, which was re-shaped into a 28 x 28 matrix and then plotted using Python.

### 5.3 Results

The partial derivative of loss function with  $L_1$  norm regularization is:

$$\frac{\partial J(w)}{\partial w} = \frac{\partial E(w)}{\partial w} + \lambda \frac{\partial |w|}{\partial w}$$

where the partial derivative of  $|w|$  can be written as:

$$\frac{\partial |w|}{\partial w_j} = \begin{cases} 1, & \text{if } w_j \geq 0 \\ -1, & \text{otherwise} \end{cases}$$

and the partial derivative in case of  $L_2$  norm regularization is:

$$\frac{\partial J(w)}{\partial w} = \frac{\partial E(w)}{\partial w} + 2\lambda w$$

The "percent correct" value was plotted over the number of training iterations for the training set, for different lambda values keeping the other hyper-parameters same.

Parameters used: **Penalty** =  $L_2$  norm, **Learning rate**  $\eta = 0.0001$




Figure 9: Accuracy Percentage v/s Epoch for  $L_2$  norm

Now for penalty as  $L_1$  norm and Learning rate  $\eta = 0.0001$ , we have:

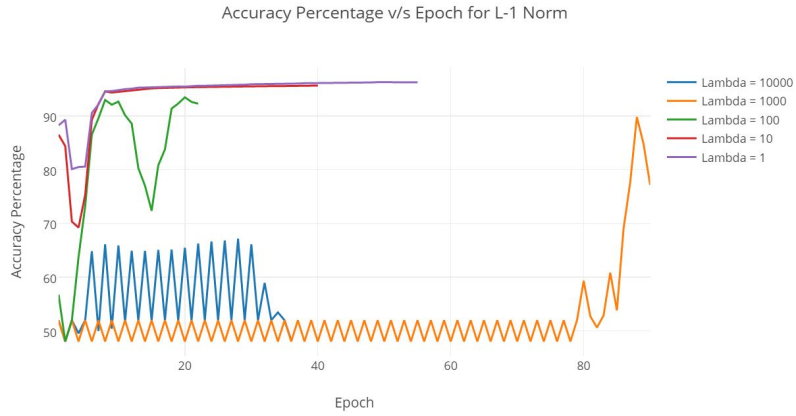


Figure 10: Accuracy Percentage v/s Epoch for  $L_1$  norm

- 153 The length of weight vector against training iterations produced a graph as follows for the  $L_2$  norm case:

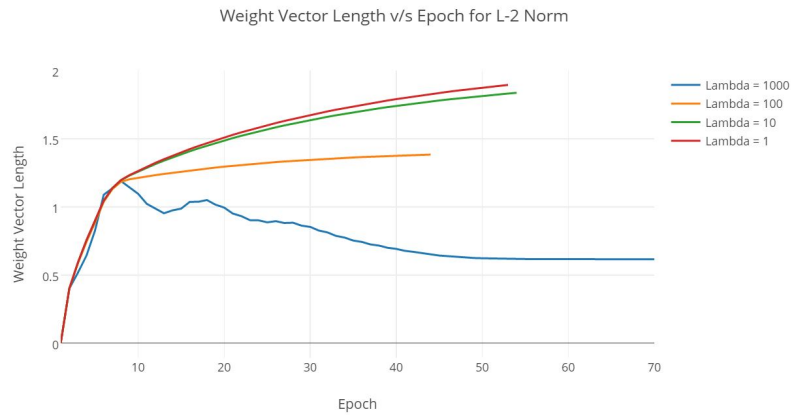


Figure 11: Weight Vector Length v/s Epoch with  $L_2$  norm

154

For  $L_1$  norm, it becomes:

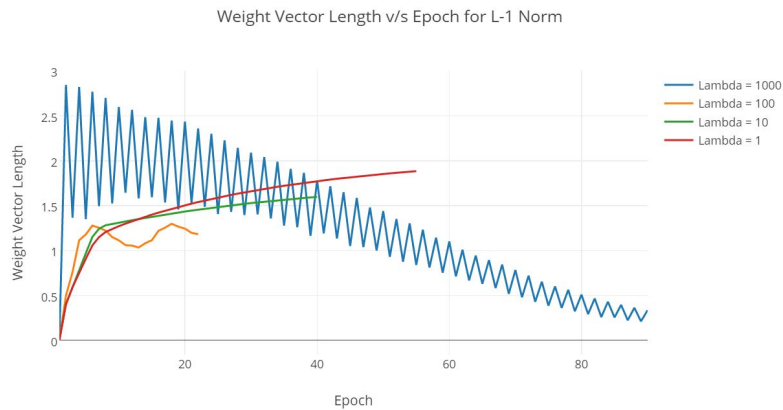


Figure 12: Weight Vector Length v/s Epoch with  $L_1$  norm

155

156 Note that the learning rate  $\eta$  used was 0.0001 in both the cases.

157 The plot of final test error for various  $\lambda$  values with  $L_2$  norm penalty is as follows:




Figure 13: Final Test Error v/s  $\log(\lambda)$   $L_2$  norm

158 The plot of final test error for various  $\lambda$  values with  $L_1$  norm penalty is as follows:




Figure 14: Final Test Error v/s  $\log(\lambda)$   $L_1$  norm

159 Note that the learning rate  $\eta$  used was 0.0001 in both the cases.

160 For **L1** norm: using learning rate  $\eta$  as 0.0001 in both the cases and different  $\lambda$  values, the final weight  
161 vectors were plotted as images. Here are the findings:

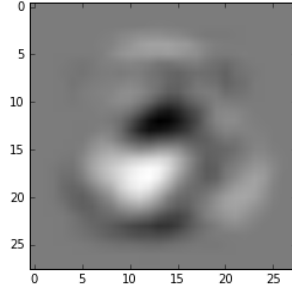


Figure 15: Weight Vector Image for 2's and 3's with  $L_1$  norm and  $\lambda = 1$

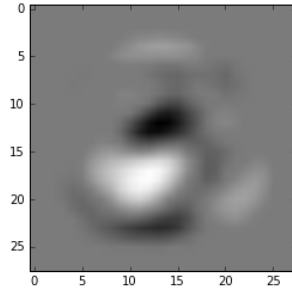


Figure 16: Weight Vector Image for 2's and 3's with  $L_1$  norm and  $\lambda = 10$

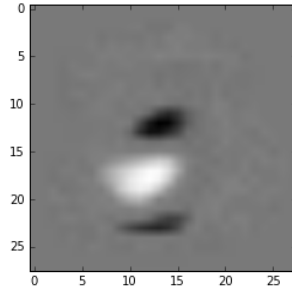


Figure 17: Weight Vector Image for 2's and 3's with  $L_1$  norm and  $\lambda = 100$

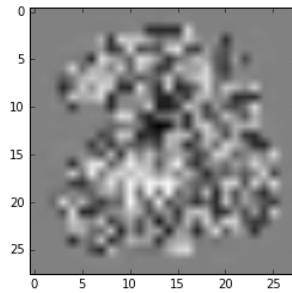


Figure 18: Weight Vector Image for 2's and 3's with  $L_1$  norm and  $\lambda = 1000$

162 For **L1** norm: Using learning rate  $\eta$  as 0.0001 in both the cases and optimal  $\lambda$  value of 0.0001, the  
 163 image looks like following:

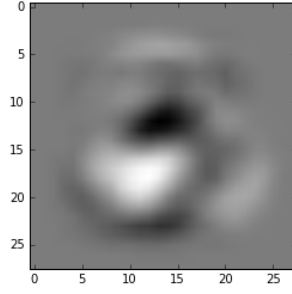


Figure 19: Weight Vector Image for 2's and 3's with  $L_1$  norm and  $\lambda = 0.0001$

164 For **L2** norm: using learning rate  $\eta$  as 0.0001 in both the cases and different  $\lambda$  values, the final weight  
 165 vectors were plotted as images. Here are the findings:

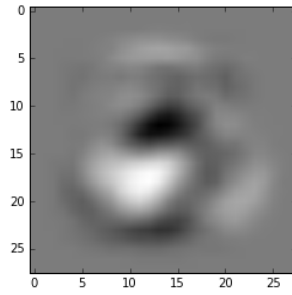


Figure 20: Weight Vector Image for 2's and 3's with  $L_2$  norm and  $\lambda = 1$

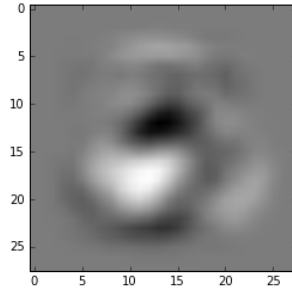


Figure 21: Weight Vector Image for 2's and 3's with  $L_2$  norm and  $\lambda = 10$

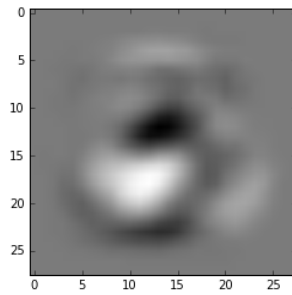


Figure 22: Weight Vector Image for 2's and 3's with  $L_2$  norm and  $\lambda = 100$




Figure 23: Weight Vector Image for 2's and 3's with  $L_2$  norm and  $\lambda = 1000$

166 For **L2** norm: Using learning rate  $\eta$  as 0.0001 in both the cases and optimal  $\lambda$  value of 0.0001, the  
 167 image looks like following:




Figure 24: Weight Vector Image for 2's and 3's with  $L_2$  norm and  $\lambda = 0.0001$

#### 168 5.4 Discussion

169 Note that the  $L_1$  norm is not differentiable at 0. However, all that matters is that we can compute  
 170 a subgradient / subderivative. Since it's differentiable everywhere else, we can just fill in any  
 171 "reasonable" value (such as -1 or 1; we have chosen -1) for the gradient at 0. So, for the cases where  
 172 the weights are negative, we would have positive regularization term that would drive them towards  
 173 zero and for the cases where the weights are positive, the regularization term would be negative, thus  
 174 decreasing the weights. Hence,  $L_1$  tries to limit the coefficient which restricts overfitting.

175 Graph ?? shows how the accuracy on training data varies according to the iteration number (epoch)  
 176 for several  $\lambda$  values.  $\lambda$  is added to avoid over-fitting. Thus, higher the  $\lambda$  value, lower will be the  
 177 accuracy on training set as more emphasis is given to the complexity function rather than the original  
 178 loss function  $E(w)$ .

179 Graph ?? shows how the accuracy on training data varies according to the iteration number (epoch)  
 180 for several  $\lambda$  values. The same generalization as above holds here as well. Note that the graph line  
 181 for  $\lambda = 10000$  overlaps with that of  $\lambda = 1000$  and is not visible clearly.

182 Next, we plot the length of weight vectors against training iterations for different  $\lambda$  values. In both the  
 183 cases, as  $\lambda$  increases, the corresponding weight vector length value decreases. Both these values are  
 184 inversely related. This is due because as  $\lambda$  increases, the overall regularized loss function value tends  
 185 to increase. However, since our goal is to minimize loss, the weight vector balances this increase in  $\lambda$   
 186 by decreasing itself. A similar argument is valid for decreasing  $\lambda$  values as well.

187 Then we have the plots of final test error for different  $\log(\lambda)$  values. In both the above graphs, as  
 188  $\lambda$  (or  $\log(\lambda)$ ) increases, the final test error increases. This is because for large  $\lambda$  values, the model  
 189 complexity is low. Beyond a certain level, the complexity may become so low that it no longer fits the  
 190 data well leading to mis-classification on a large number of points. Increasing  $\lambda$  value only decreases  
 191 the complexity further, leading to even further decline in test accuracy. Similarly, if the  $\lambda$  value is too  
 192 low, the model may become highly complex, so much so that over-fitting happens. Test error in this  
 193 case will again be large, as the model won't generalize well for points in test set.



194 Finally, we have the weight vector images. Since, regularization limits the weights from being too  
195 high, the images of weights are little bit softer in nature.

## 6 Softmax Regression Via Gradient Descent

### 6.1 Introduction

The task at hand is to perform Softmax Regression on the MNIST data set and come up with the best parameters that may perform well on the test data, without actually looking at test data. Then, we need to plot the loss function values over number of training iterations for training, hold-out and test data sets. Finally, we are expected to plot the percent correct values over training iterations the three data set parts.

### 6.2 Methodology

Softmax Regression was performed on the first 20,000 training data points and was used to do a 10-way classification of the hand-written digits using a hold-out set and regularization parameter  $\lambda$ . The size of hold-out set was again set to 10% of the size of training data. To figure out the best hyper-parameter values, a hold-out set was used. The parameters performing the best on this hold-out set were chosen to be the final parameters. The loss function(E) was plotted over the number of training iterations for the training, hold-out and test sets.

Parameters used: Penalty =  $L_2$  norm; Regularization parameter  $\lambda = 0.0001$ ; Learning rate  $\eta = 0.0001$

Next, the "percent correct" (or accuracy percentage) was plotted over the number of training iterations for the training, hold-out and test sets. Same parameters as above were used.

### 6.3 Results

The percentage error values recorded on the hold-out set for different values of hyper-parameters are as follows:

Table 1: Error on Hold-out Set for Different Hyper-parameters

$\eta$	$\lambda$	Norm	Error %
0.0001	0.01	L-1/L-2	8.1
0.0001	0.1	L-1/L-2	8.1
0.0001	0.0	L-1/L-2	8.1
0.001	0.0	L-1/L-2	8.15
0.01	0.0	L-1/L-2	8.2
0.1	0.0	L-1/L-2	88.2

All the graphs produced are plotted and reported herein.

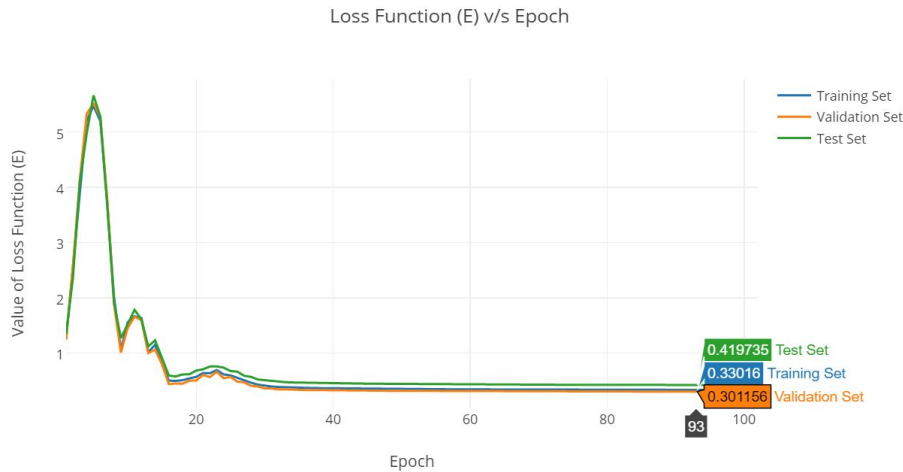


Figure 25: Loss Function(E) v/s Epoch

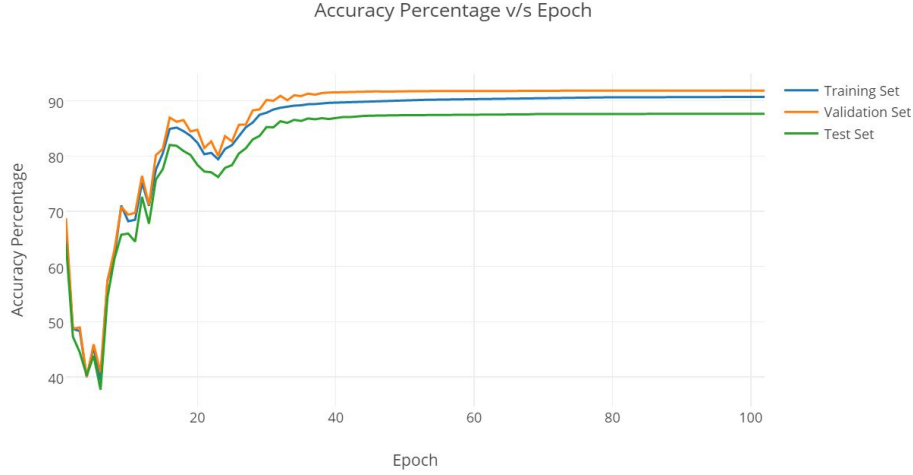


Figure 26: Accuracy Percentage v/s Epoch

## 6.4 Discussion

As evident from the given table, the optimum value of error on hold-out (validation) set was obtained at more than one pair of values of  $\eta$  and  $\lambda$ . We decided to go with the highest  $\eta$  and the largest non-zero  $\lambda$  that gave us the optimum error value on hold-out set. This is because we wanted  $\lambda$  as large as possible to help generalization and avoid over-fitting, and at the same time being a good representative of the data. Also, higher learning rate was preferred for faster convergence. Since the penalties at both  $L_2$  and  $L_1$  norm did not seem to affect the error values, we decided to go forward with  $L_2$  norm, as it gives a better measure of loss and is convex everywhere.

Thus, the hyper-parameters chosen were": Penalty =  $L_2$  norm Regularization parameter  $\lambda = 0.1$   
Learning rate  $\eta = 0.0001$

Validation Error = 8.1

Test Error obtained = 12.35

Note that early stopping was used to make sure that we do not over-fit the data.

Graph ?? shows how the training, validation and test errors vary according to the iteration number (epoch) for the same values of hyper-parameters. As evident from the graph, the loss function decreases and stabilizes over time, which corresponds to convergence. As an indicative measure, the values of test, validation and training errors have been displayed for a particular value (93) of epoch.

Graph ?? shows how the training, validation and test percent correct values vary according to the iteration number (epoch) for the same values of hyper-parameters. As evident from the graph, the accuracy saturates over time and does not improve. The weights learned give the best performance on validation set, followed by training and test set. It is interesting to note that all the values are fairly close to each other and similar in shape, which means that the training set is a good representative of the hold-out and test sets.

## 7 Results and Learnings

The best accuracy achieved using Logistic Regression was more than 97% for both the subsets of data - having digits 2 and 3, and having digits 2 and 8. The best learning rate was  $\eta = 0.0001$  and the best regularization parameter  $\lambda$  was 0.0001 for 2's and 3's, and 0.1 for 2's and 8's. Using Softmax Regression, our accuracy was around 87.65% on test data. The  $\lambda$  used in this case was 0.001 while the learning rate was kept to be 0.0001. Annealing of learning rate per iteration helped us avoid over-fitting of the data, even when number of iterations was huge. This annealing parameter was set to a convenient value (2000) to ensure a gradual yet sufficient decrease in learning rate. The early stopping margin for number of iterations was set to 15 for the both the cases to ensure that the algorithm was not stopping at some sub-optimal value, which was the case when this value was small.

Having taken courses like 250A and 250B, we knew the working and mechanism of Logistic and Softmax Regression, but we had never had a chance to perform their in-depth analysis by ourselves. Deriving the expressions and implementing these two algorithms along with regularization gave us a new insight and deep understanding of the working of these methods. Questions involving plotting of loss functions, weight vectors as images and weight vector lengths showed us how these values vary with different regularization parameters  $\lambda$  and the iteration number. We had never looked at the MNIST digit classification problem from this perspective and now have a clearer idea as to how the various hyper-parameters are related to each other. The impact of regularization on the model and weights is now in front of us, while previously we only looked at it theoretically. The concepts of annealing and early stopping were entirely new to us, as these concepts were never visible when using *scikit* for performing these computations.

## 8 Individual Contributions

Being roommates, it was extremely convenient and simple for both the authors to coordinate and work in sync, while ensuring equal distribution of work and time spent on the assignment. Whenever one of the authors got stuck at some point, the other was there to help him out and unblock instantly. The process was initially started with both of us sitting down and solving on a white board the various derivations involved. The work thereafter was taken up as under:

Chetan Gandotra implemented the part which involved reading of MNIST data and implemented Softmax Regression. Then, debugging and graph plotting of Logistic Regression was taken up by him.

Rishabh Misra took up the implementation of Logistic Regression after extracting digit specific data. Thereafter, debugging and graph plotting of Softmax Regression was taken up by him.

The implementation and graphs for regularization question (Q5) were divided equally, with 5 (b) and 5 (c) being taken up by Chetan, and the others by Rishabh. For parameter tuning (values of  $\lambda$ ,  $\eta$ , T, early stopping iteration number etc.), we made an Excel sheet with possible list of values and divided them equally amongst us. We then ran the code for these values for parameters on our respective systems. When it came to writing the report, we took alternate question parts, with Rishabh taking odd questions and Chetan taking up even questions.

## 278 References

279 [1] <https://github.com/akosiorek/CSE/tree/master/MLCV>

## 280 Appendix

### 281 LoadMNIST.py

```
282 import os, struct
283 from array import array as pyarray
284 from numpy import append, array, int8, uint8, zeros
285
286 def load_mnist(dataset="training", digits=None, path=None, asbytes=False, selection=None,
287 return_labels=True, return_indices=False):
```

```
288     """
289     Loads MNIST files into a 3D numpy array.
```

```
290
291     You have to download the data separately from [MNIST]_. It is recommended
292     to set the environment variable ‘‘MNIST’’ to point to the folder where you
293     put the data, so that you don’t have to select path. On a Linux+bash setup,
294     this is done by adding the following to your ‘‘.bashrc’’::
```

```
295
296         export MNIST=/path/to/mnist
```

### 297 Parameters

```
298
299
300 dataset : str
301     Either "training" or "testing", depending on which dataset you want to
302     load.
303 digits : list
304     Integer list of digits to load. The entire database is loaded if set to
305     ‘‘None’’. Default is ‘‘None’’.
306 path : str
307     Path to your MNIST datafiles. The default is ‘‘None’’, which will try
308     to take the path from your environment variable ‘‘MNIST’’. The data can
309     be downloaded from http://yann.lecun.com/exdb/mnist/.
310 asbytes : bool
311     If True, returns data as ‘‘numpy.uint8’’ in [0, 255] as opposed to
312     ‘‘numpy.float64’’ in [0.0, 1.0].
313 selection : slice
314     Using a ‘slice’ object, specify what subset of the dataset to load. An
315     example is ‘‘slice(0, 20, 2)’’, which would load every other digit
316     until—but not including—the twentieth.
317 return_labels : bool
318     Specify whether or not labels should be returned. This is also a speed
319     performance if digits are not specified, since then the labels file
320     does not need to be read at all.
321 return_indices : bool
322     Specify whether or not to return the MNIST indices that were fetched.
323     This is valuable only if digits is specified, because in that case it
324     can be valuable to know how far
325     in the database it reached.
```

### 326 Returns

```
327
328
329 images : ndarray
330     Image data of shape ‘‘(N, rows, cols)’’, where ‘‘N’’ is the number of images. If
331     neither labels nor indices are returned, then this is returned directly, and not inside
332 labels : ndarray
333     Array of size ‘‘N’’ describing the labels. Returned only if ‘‘return_labels’’ is
334     ‘‘True’’, which is default.
335 indices : ndarray
336     The indices in the database that were returned.
```

337

## Examples

Assuming that you have downloaded the MNIST database and set the environment variable ‘\$MNIST’ point to the folder, this will load all images and labels from the training set:

```
>>> images, labels = ag.io.load_mnist('training') # doctest: +SKIP
```

Load 100 sevens from the testing set:

```
>>> sevens = ag.io.load_mnist('testing', digits=[7], selection=slice(0, 100),
    return_labels=False) # doctest: +SKIP
```

```
"""
```

```
# The files are assumed to have these names and should be found in 'path'
```

```
files = {
    'training': ('train-images.idx3-ubyte', 'train-labels.idx1-ubyte'),
    'testing': ('t10k-images.idx3-ubyte', 't10k-labels.idx1-ubyte'),
}
```

```
if path is None:
```

```
    try:
```

```
        path = 'C:\\Users\\Chetan\\Documents\\Python Scripts\\Way1'
```

```
        #path = os.environ['MNIST']
```

```
    except KeyError:
```

```
        raise ValueError("Unspecified path requires environment variable $MNIST
            to be set")
```

```
try:
```

```
    images_fname = os.path.join(path, files[dataset][0])
```

```
    labels_fname = os.path.join(path, files[dataset][1])
```

```
except KeyError:
```

```
    raise ValueError("Data set must be 'testing' or 'training'")
```

```
# We can skip the labels file only if digits aren't specified and labels aren't
asked for
```

```
if return_labels or digits is not None:
```

```
    flbl = open(labels_fname, 'rb')
```

```
    magic_nr, size = struct.unpack(">II", flbl.read(8))
```

```
    labels_raw = pyarray("b", flbl.read())
```

```
    flbl.close()
```

```
fimg = open(images_fname, 'rb')
```

```
magic_nr, size, rows, cols = struct.unpack(">IIII", fimg.read(16))
```

```
images_raw = pyarray("B", fimg.read())
```

```
fimg.close()
```

```
if digits:
```

```
    indices = [k for k in range(size) if labels_raw[k] in digits]
```

```
else:
```

```
    indices = range(size)
```

```
if selection:
```

```
    indices = indices[selection]
```

```
N = len(indices)
```

```
images = zeros((N, rows, cols), dtype=uint8)
```

```
if return_labels:
```

```
    labels = zeros((N), dtype=int8)
```

```
for i, index in enumerate(indices):
```

```
    images[i] = array(images_raw[ indices[i]*rows*cols : (indices[i]+1)*rows*cols
        ]).reshape((rows, cols))
```

```
    if return_labels:
```

```

403         labels[i] = labels_raw[indices[i]]
404
405     if not asbytes:
406         images = images.astype(float)/255.0
407
408     ret = (images,)
409     if return_labels:
410         ret += (labels,)
411     if return_indices:
412         ret += (indices,)
413     if len(ret) == 1:
414         return ret[0] # Don't return a tuple of one
415     else:
416         return ret
417
418 Logistic_Regression_via_Gradient_Descent.py
419
420 """
421 CSE 253: Neural Networks and Pattern Recognition
422 Logistic Regression With and Without Gradient Descent
423 This file contains code for questions 4 and 5, including all graph plots
424 """
425 import numpy
426 import math
427 import matplotlib.pyplot as plt
428 import plotly.plotly as pyl
429 import plotly.graph_objs as go
430
431 from LoadMNIST import load_mnist
432 #-----Utility functions-----
433 def get_data(N, N_test):
434     #load MNIST data using libraries available
435     training_data, training_labels = load_mnist('training')
436     test_data, test_labels = load_mnist('testing')
437
438     training_data = flatArray(N, 784, training_data) #training_data is N x 784 matrix
439     training_labels = training_labels[:N]
440     test_data = flatArray(N_test, 784, test_data)
441     test_labels = test_labels[:N_test]
442
443     # adding column of 1s for bias
444     training_data = addOnesColAtStart(training_data)
445     test_data = addOnesColAtStart(test_data)
446
447     # Last 10% of training data size will be considered as the validation set
448     N_validation = int(N / 10)
449     validation_data = training_data[N-N_validation:N]
450     validation_labels = training_labels[N-N_validation:N]
451     N=N-N_validation
452     #update training data to remove validation data
453     training_data = training_data[:N]
454     training_labels = training_labels[:N]
455
456     return training_data, training_labels, test_data, test_labels, validation_data,
457         validation_labels
458
459 def flatArray(rows, cols, twoDArr):
460     flattened_arr = numpy.zeros(shape=(rows, cols))
461     for row in range(0, rows):
462         i=0
463         for element in twoDArr[row]:
464             for ell in element:
465                 flattened_arr[row][i] = ell
466                 i = i+1
467     return flattened_arr

```

```

468
469 def addOnesColAtStart(matrix):
470     Ones = numpy.ones(len(matrix))
471     newMatrix = numpy.c_[Ones, matrix]
472     return newMatrix
473
474 # custom sigmoid function; if -x is too large, return value 0
475 def sigmoid(x):
476     if(-x < 709):
477         return 1 / (1 + math.exp(-x))
478     else:
479         return 1 / (1 + math.exp(708))
480
481 def extract_digit_specific_data(digits, data, label):
482     pruned_data = numpy.zeros(shape = (1, len(data[0])))
483     pruned_labels = []
484     cnt = 0
485     for i in range(0, len(label)):
486         if label[i] in digits:
487             if (cnt == 0):
488                 for j in range(len(data[0])):
489                     pruned_data[0][j] = data[i][j]
490             else:
491                 pruned_data = addRowToMatrix(pruned_data, data[i])
492                 pruned_labels.append(digits.get(label[i]))
493                 cnt = cnt + 1
494     return pruned_data, pruned_labels
495
496 def addRowToMatrix(matrix, row):
497     newMatrix = numpy.zeros(shape = ((len(matrix) + 1), len(matrix[0])))
498     for i in range(len(matrix)):
499         newMatrix[i] = matrix[i]
500     newMatrix[i+1] = row
501     return newMatrix
502
503 def calculate_log_likelihood(data, label, weights, t):
504     log_likelihood = 0.0
505     for j in range(0,t):
506         #print(numpy.log(sigmoid(numpy.dot(weights, data[j]))))
507         log_likelihood += (label[j]*numpy.log(sigmoid(numpy.dot(weights, data[j])))) +
508             ((1-label[j])*numpy.log(sigmoid(-1*numpy.dot(weights, data[j]))))
509
510     return -1*log_likelihood/t
511
512 # 5.b, 5.c – Plot of Percent correct on training data v/s iterations,
513 #length of weight vector v/s lambda
514 def plotlyGraphsRegularization(error_plot_array, lamda_vals, graph_name, min_index = -1):
515     pyl.sign_in('chetang', 'vil7vTAuCSWt2IEZvaH9')
516
517     trace = []
518
519     for i in range(len(lamda_vals)):
520         y1 = error_plot_array[i]
521         #y1 = y1[:min_index[i]]
522         x1 = [j+1 for j in range(len(y1))]
523
524         trancel = go.Scatter(
525             x=x1,
526             y=y1,
527             name = 'lambda = ' + (str)(lamda_vals[i]), # Style name/legend entry with
528             html tags
529             connectgaps=True
530         )
531
532         trace.append(trancel)

```



```

533     data = trace
534
535     fig = dict(data=data)
536     pyl.ipplot(fig, filename=graph_name)
537
538 def plotlyGraphs(error_plot_array, labels, name):
539     pyl.sign_in('cgandotr', '3c9fho4498')
540     trace = []
541
542     for i in range(len(labels)):
543         y1 = error_plot_array[i]
544         y1 = [k for k in y1]
545         x1 = [(j+1) for j in range(len(y1))]
546
547         trace1 = go.Scatter(
548             x=x1,
549             y=y1,
550             name = str(labels[i]), # Style name/legend entry with html tags
551             connectgaps=True
552         )
553         trace.append(trace1)
554     data = trace
555     fig = dict(data=data)
556     pyl.ipplot(fig, filename=name)
557
558 def early_stopping(early_stopping_horizon, accuracy, i):
559     if i > early_stopping_horizon + 1:
560         counter = 0;
561         for p in range(0, early_stopping_horizon):
562             if accuracy[i-p] <= accuracy[i-p-1]:
563                 counter+=1
564             else:
565                 break
566
567         if(counter == early_stopping_horizon):
568             return True
569         else:
570             return False
571
572 def calculate_error(weights, data, label):
573     error = 0.0;
574     for j in range(0, len(data)):
575         prediction = sigmoid(numpy.dot(weights, data[j]))
576         if(prediction > 0.5 and label[j] != 1):
577             error += 1;
578         elif (prediction <= 0.5 and label[j] != 0):
579             error += 1;
580     return error;
581
582 def dropFirstColumn(weights):
583     return numpy.array(weights)[0][1:]
584
585 def fit(training_data, training_label, test_data, test_label, validation_data,
586         validation_label, digits, learning_rate=0.0001, iteration=200, batch_size=0,
587         T=5000, lamda_vals=[0], norm=2):
588
589     t = len(training_data)
590     accuracy_plot_array = []
591     log_likelihood_array = []
592     weight_vector_length_array = []
593     accuracy_plot_training_array = []
594     weights_for_all_lamda = []
595     test_error_array = []
596     org_learning_rate = learning_rate
597

```

```

598     for lamda in lamda_vals:
599         weights = numpy.matrix(numpy.zeros(len(training_data[0])))
600         weights_array = []
601         weight_vector_length = []
602
603         accuracy_plot_training = []
604         accuracy_plot_validation = []
605         accuracy_plot_testing = []
606
607         log_likelihood_training = []
608         log_likelihood_validation = []
609         log_likelihood_testing = []
610         min_error_index = 0
611         learning_rate = org_learning_rate
612
613     for i in range(0, iteration):
614         # initialise Gradient
615         gradient = numpy.matrix(numpy.zeros(len(training_data[0])))
616         cnt = 0
617
618         norm_term = []
619         if (norm == 2):
620             norm_term = l2_norm(lamda, weights)
621         else:
622             norm_term = l1_norm(lamda, weights)
623
624         # calculate gradient over all the samples
625         for j in range(0, t):
626             # update gradient
627             gradient += ((training_label[j]) - sigmoid(numpy.dot(weights,
628             training_data[j]))) * numpy.array(training_data[j])
629             cnt += 1
630             if (batch_size != 0 and cnt == (int)(t/batch_size)):
631                 cnt = 0
632                 weight_vector_length.append(numpy.linalg.norm(weights))
633                 weights = weights + learning_rate * (gradient - norm_term)
634                 # re-initialise Gradient
635                 gradient = numpy.matrix(numpy.zeros(len(training_data[0])))
636                 # calculating log likelihood of training, validation and test
637                 dataset
638                 log_likelihood_training.append(calculate_log_likelihood(training_data,
639                 training_label, weights, t))
640                 log_likelihood_validation.append(calculate_log_likelihood(validation_data,
641                 validation_label, weights, len(validation_data)))
642                 log_likelihood_testing.append(calculate_log_likelihood(test_data,
643                 test_label, weights, len(test_data)))
644
645             if (batch_size == 0):
646                 weight_vector_length.append(numpy.linalg.norm(weights))
647                 # update weights vector according to the update rule of Gradient descent method
648                 weights = weights + learning_rate * (gradient - norm_term)
649
650             # calculating log likelihood of training, validation and test dataset
651             log_likelihood_training.append(calculate_log_likelihood(training_data,
652             training_label, weights, t))
653             log_likelihood_validation.append(calculate_log_likelihood(validation_data,
654             validation_label, weights, len(validation_data)))
655             log_likelihood_testing.append(calculate_log_likelihood(test_data, test_label,
656             weights, len(test_data)))
657
658         # annealing of learning rate
659         learning_rate = learning_rate/(1+i/T)
660
661         # calculating error percentage on train, test and validation data
662         accuracy_plot_training.append((len(training_data) - calculate_error(weights,

```

```

663         training_data , training_label))*100/len(training_data))
664         accuracy_plot_validation.append((len(validation_data) - calculate_error(weights ,
665         validation_data , validation_label))*100/len(validation_data))
666         accuracy_plot_testing.append((len(test_data) - calculate_error(weights , test_data ,
667         test_label))*100/len(test_data))
668
669         weights_array.append(weights)
670
671         # check for early stopping
672         early_stopping_horizon = 15
673         min_error_index = i
674         if (early_stopping(early_stopping_horizon , accuracy_plot_validation , i) and i >
675         early_stopping_horizon):
676             min_error_index = i-early_stopping_horizon;
677             weights = weights_array[min_error_index]
678             break
679
680         weight_vector_length_array.append(weight_vector_length)
681         weights_for_all_lambda.append(weights)
682
683         log_likelihood_array.append(log_likelihood_training);
684         log_likelihood_array.append(log_likelihood_validation);
685         log_likelihood_array.append(log_likelihood_testing);
686         accuracy_plot_array.append(accuracy_plot_training)
687         accuracy_plot_array.append(accuracy_plot_validation)
688         accuracy_plot_array.append(accuracy_plot_testing)
689         accuracy_plot_training_array.append(accuracy_plot_training)
690
691         test_error = (calculate_error(weights , test_data , test_label)*100)/len(test_data)
692         validation_error = (calculate_error(weights , validation_data , validation_label)*100)/
693         len(validation_data)
694         print('Error on validation dataset : ' + str(validation_error) + '%');
695         print('Error on test dataset : ' + str(test_error) + '%');
696         test_error_array.append(test_error)
697
698         #For single lambda value
699         if (len(lamda_vals) == 1):
700             plotlyGraphs(log_likelihood_array , ['Training Set','Validation Set','Test Set'],
701             'Log Likelihood Plot')
702             plotlyGraphs(accuracy_plot_array , ['Training Set','Validation Set','Test Set'],
703             'Percent Correct Plot')
704         #else:
705             #For multiple lambda values
706             #plotlyGraphsRegularization(accuracy_plot_training_array , lamda_vals , "Accuracy vs
707             Epoch")
708             #plotlyGraphsRegularization(weight_vector_length_array , lamda_vals , "Weight Vector
709             Length vs Epoch")
710             #plotlyErrorVsLamda(test_error_array , lamda_vals)
711
712         # printing error on training and testing dataset
713         return weights_for_all_lambda
714
715     def plot_weights(weights):
716         lamda_vals = [1000, 100, 10, 1, 0.1, 0.001, 0.0001]
717         i = 0
718         for w in weights:
719             # Plot weights as image after removing bias terms. The rest of columns are pixels
720             print ('lambda = ' + str(lamda_vals[i]))
721             i += 1
722             pixels1 = dropFirstColumn(w)
723             pixels = numpy.reshape(pixels1 , (28, 28))
724             plt.imshow(pixels , cmap='gray')
725             plt.show()
726
727     def l2_norm(lamda , weights):

```

```

728     return 2*lamda*weights;
729
730 def ll_norm(lamda, weights):
731     w = numpy.ones(len(weights))
732     for i in range(len(weights)):
733         if (weights[i] < 0):
734             w[i] = -1
735     return lamda*w;
736
737 # 5.d – Final test error v/s Lambda values graph
738 # Generates bar graphs
739 def plotlyErrorVsLambda(test_error_array, lamda_vals):
740     lamda_vals1 = [math.log(lamda) for lamda in lamda_vals]
741     pyl.sign_in('chetang', 'vil7vTAuCSWt2lEZvaH9')
742     trace = []
743     colors = ['rgb(104,224,204)', 'rgb(204,204,204)', 'rgb(49,130,189)', 'rgb(41,180,129)']
744     for i in range(0, len(test_error_array)):
745         y1 = test_error_array[i]
746         x1 = lamda_vals1[i]
747
748         trace1 = go.Bar(
749             x=x1,
750             y=y1,
751             name='Lambda = log(' + (str)(lamda_vals[i]) + ')',
752             marker=dict(
753                 color=colors[i]
754             )
755         )
756         trace.append(trace1)
757
758     layout = go.Layout(
759         xaxis=dict(tickangle=-45),
760         barmode='group',
761     )
762     fig = go.Figure(data=trace, layout=layout)
763     pyl.iplot(fig, filename='Final Error vs Lambda')
764
765 # -----Main function-----
766
767 if __name__ == "__main__":
768     numpy.random.seed(0)
769
770     N = 20000
771     N_test = 2000
772     iteration = 200
773     batch_size = 0
774     T = 2000
775
776     #lamda_vals = [0]
777     lamda_vals = [1000, 100, 10, 1, 0.1, 0.001, 0.0001] # regularization
778     weightage parameter.
779     Put [0] for no regularization
780     norm = 2 # 2 for l-2 norm, 1 for l-1 norm
781
782     full_training_data, full_training_label, full_test_data, full_test_label,
783     full_validation_data, full_validation_label = get_data(N, N_test)
784
785     # Parameters for training data on 2's and 3's
786     digits = {2:1, 3:0}
787     training_data, training_label = extract_digit_specific_data(digits,
788         full_training_data,
789         full_training_label)
790     validation_data, validation_label = extract_digit_specific_data(digits,
791         full_validation_data,
792         full_validation_label)

```

```

793     test_data , test_label = extract_digit_specific_data(digits , full_test_data ,
794     full_test_label)
795     learning_rate = 0.0001
796
797     weights23 = fit(training_data , training_label , test_data , test_label ,
798                     validation_data , validation_label , learning_rate , iteration ,
799                     batch_size , digits , T, lamda_vals , norm)
800     plot_weights(weights23)
801
802     # Parameters for training data on 2's and 8's
803     digits = {2:1, 8:0}
804     training_data , training_label = extract_digit_specific_data(digits ,
805     full_training_data ,
806     full_training_label)
807     validation_data , validation_label = extract_digit_specific_data(digits ,
808     full_validation_data ,
809     full_validation_label)
810     test_data , test_label = extract_digit_specific_data(digits , full_test_data ,
811     full_test_label)
812     learning_rate = 0.1
813     lamda_vals = [0] #Regularization not asked for with 2/8 case
814
815     weights28 = fit(training_data , training_label , test_data , test_label ,
816                     validation_data , validation_label , digits , learning_rate ,
817                     iteration ,
818                     batch_size , T, lamda_vals , norm)
819     plot_weights(weights28)
820
821     weights = weights28[0] - weights23[0]
822
823     plot_weights(weights)
824
825
826
827 Softmax_Regression.py
828
829 """
830 Softmax Regression on MNIST Data set to perform 10-way classification
831 """
832 import numpy
833 import math
834 import plotly.plotly as pyl
835 import plotly.graph_objs as go
836
837 from LoadMNIST import load_mnist
838
839 #-----Utility functions-----
840 def get_data(N, N_test):
841     #load MNIST data using libraries available
842     training_data , training_labels = load_mnist('training ')
843     test_data , test_labels = load_mnist('testing ')
844
845     training_data = flatArray(N, 784, training_data) #training_data is N x
846     784 matrix
847     training_labels = training_labels[:N]
848     test_data = flatArray(N_test, 784, test_data)
849     test_labels = test_labels[:N_test]
850
851     # adding column of 1s for bias
852     training_data = addOnesColAtStart(training_data)
853     test_data = addOnesColAtStart(test_data)
854
855     # Last 10% of training data size will be considered as the validation
856     set
857     N_validation = int (N / 10)

```

```

858     validation_data = training_data[N-N_validation:N]
859     validation_labels = training_labels[N-N_validation:N]
860     N=N-N_validation
861     #update training data to remove validation data
862     training_data = training_data[:N]
863     training_labels = training_labels[:N]
864
865     return training_data , training_labels , test_data , test_labels ,
866           validation_data ,
867           validation_labels
868
869 def flatArray(rows , cols , twoDArr):
870     flattened_arr = numpy.zeros(shape=(rows , cols))
871     for row in range(0 , rows):
872         i=0
873         for element in twoDArr[row]:
874             for ell in element:
875                 flattened_arr[row][i] = ell
876                 i = i+1
877     return flattened_arr
878
879 def addOnesColAtStart(matrix):
880     Ones = numpy.ones(len(matrix))
881     newMatrix = numpy.c_[Ones , matrix]
882     return newMatrix
883
884 # custom sigmoid function; if -x is too large , return value 0
885 def exp(x):
886     if(x < 709):
887         return math.exp(x)
888     else:
889         return math.exp(709)
890
891 def addRowToMatrix(matrix , row):
892     newMatrix = numpy.zeros(shape = ((len(matrix) + 1), len(matrix[0])))
893     for i in range(len(matrix)):
894         newMatrix[i] = matrix[i]
895     newMatrix[i+1] = row
896     return newMatrix
897
898 def l2_norm(lamda , weights):
899     return 2*lamda*weights;
900
901 def l1_norm(lamda , weights):
902     w = numpy.ones((len(numpy.array(weights)),
903                     len(numpy.array(weights)[0])))
904     for i in range(len(w)):
905         for j in range(len(w[0])):
906             if (weights[i,j] < 0):
907                 w[i][j] = -1
908     return lamda*numpy.matrix(w);
909
910 # custom sigmoid function; if -x is too large , return value 0
911 def sigmoid(x):
912     if(-x < 709):
913         return 1 / (1 + math.exp(-x))
914     else:
915         return 1 / (1 + math.exp(708))
916
917 def calculate_log_likelihood(data , label , weights , t):
918     log_likelihood = 0.0
919     for j in range(0,t):
920         log_likelihood += (label[j]*numpy.log(sigmoid(numpy.dot(weights ,
921                                                                 data[j]))) +
922                             ((1-label[j])*numpy.log(sigmoid(-1*numpy.dot(weights , data[j])))))

```

```

923
924     return -l*log_likelihood/t
925
926 def plotlyGraphs(error_plot_array, labels, name):
927     pyl.sign_in('cgandotr', '3c9fho4498')
928     trace = []
929     for i in range(len(labels)):
930         y1 = error_plot_array[i]
931         y1 = [k for k in y1]
932         #y1 = y1[:min_index[i]]
933         x1 = [(j+1) for j in range(len(y1))]
934
935         trace1 = go.Scatter(
936             x=x1,
937             y=y1,
938             name = str(labels[i]), # Style name/legend entry with html tags
939             connectgaps=True
940         )
941
942         trace.append(trace1)
943     data = trace
944     fig = dict(data=data)
945     pyl.iplot(fig, filename=name)
946
947 def early_stopping(early_stopping_horizon, accuracy, i):
948     if i > early_stopping_horizon:
949         counter = 0;
950         for p in range(0, early_stopping_horizon):
951             if accuracy[i-p] <= accuracy[i-p-1]:
952                 counter+=1
953             else:
954                 break
955
956         if(counter == early_stopping_horizon):
957             return True
958         else:
959             return False
960
961 def calculate_error(weights, data, label, k = 10):
962     error = 0.0;
963     for j in range(0, len(data)):
964         softmax_denom = 0.0
965         softmax_num = numpy.zeros(k);
966         for x in range(0, k):
967             softmax_num[x] = exp(numpy.dot(numpy.transpose(weights[:, x]),
968                 data[j]))
969             softmax_denom += softmax_num[x]
970
971         prediction = numpy.argmax(softmax_num/softmax_denom)
972         if(prediction != label[j]):
973             error += 1;
974     return error;
975
976 def softmax_loss(weights, labels, data, c):
977     loss = 0.0
978     for i in range(len(data)):
979         softmax_denom = 0.0
980         softmax_num = numpy.zeros(c);
981         for x in range(0, c):
982             softmax_num[x] = exp(numpy.dot(numpy.transpose(weights[:, x]),
983                 data[i]))
984             softmax_denom += softmax_num[x]
985         softmax = softmax_num[labels[i]]/softmax_denom
986         if (softmax > 0):
987             log_softmax = math.log(softmax)

```

```

988         else:
989             log_softmax = 0
990             loss += (log_softmax)
991         return -1*loss/(len(data))
992
993 def getSoftmax(k, weights, training_data, j):
994     softmax_denom = 0.0
995     softmax_num = numpy.zeros(k);
996     for x in range(0,k):
997         softmax_num[x] = exp(numpy.dot(numpy.transpose(weights[:,x]),
998             training_data[j]))
999         softmax_denom += softmax_num[x]
1000     return softmax_num/softmax_denom
1001
1002 def fit(training_data, training_label, test_data, test_label, validation_data,
1003     validation_label, iteration = 1000, T=2000, lamda=0.001,
1004     learning_rate = 0.0001, norm = 2):
1005     k = len(numpy.unique(training_label))
1006     t = len(training_data)
1007     weights = numpy.matrix(numpy.zeros((len(training_data[0]), k)))
1008     weights_array = []
1009
1010     early_stopping_horizon = 15
1011     error_plot = numpy.zeros(iteration)
1012     min_error_index = 0
1013     loss_array = []
1014     loss_training = []
1015     loss_validation = []
1016     loss_testing = []
1017
1018     accuracy_plot_array = []
1019     accuracy_plot_training = []
1020     accuracy_plot_validation = []
1021     accuracy_plot_testing = []
1022     test_error = 0.0
1023
1024     for i in range(0, iteration):
1025         # initialise Gradient
1026         gradient = numpy.matrix(numpy.zeros((len(training_data[0]), k)))
1027
1028         # calculate gradient over all the samples
1029         for j in range(0,t):
1030             modified_label = numpy.zeros(k);
1031             modified_label[training_label[j]] = 1
1032             softmax = getSoftmax(k, weights, training_data, j)
1033             gradient += numpy.transpose(numpy.transpose(numpy.matrix(
1034                 modified_label - softmax)
1035                 ) * numpy.matrix(training_data[j]))
1036
1037         norm_term = 0.0
1038         if (norm == 2):
1039             norm_term = l2_norm(lamda, weights)
1040         else:
1041             norm_term = l1_norm(lamda, weights)
1042         # update weights vector according to the update rule of Gradient
1043         # descent method
1044         weights = weights + learning_rate * (gradient - norm_term)
1045
1046         loss_training.append(softmax_loss(weights, training_label,
1047             training_data, k))
1048         loss_validation.append(softmax_loss(weights, validation_label,
1049             validation_data, k))
1050         loss_testing.append(softmax_loss(weights, test_label, test_data, k))
1051
1052         # calculating error percentage on train, test and validation data

```



```

1053     training_error = calculate_error(weights, training_data,
1054     training_label)
1055     validation_error = calculate_error(weights, validation_data,
1056     validation_label)
1057     test_error = calculate_error(weights, test_data, test_label)
1058
1059     accuracy_plot_training.append((len(training_data) - training_error)
1060     *100/len(training_data))
1061     accuracy_plot_validation.append((len(validation_data) -
1062     validation_error)*100/len(validation_data))
1063     accuracy_plot_testing.append((len(test_data) - test_error)*100/
1064     len(test_data))
1065
1066     learning_rate = learning_rate/(1+i/T)
1067
1068     error_plot[i] = validation_error*100/len(validation_data);
1069     weights_array.append(weights)
1070
1071     # check for early stopping
1072     if (early_stopping(early_stopping_horizon, accuracy_plot_validation,
1073     i)):
1074         min_error_index = i-early_stopping_horizon;
1075         weights = weights_array[min_error_index]
1076         break
1077     min_error_index = i
1078
1079     loss_array.append(loss_training)
1080     loss_array.append(loss_validation)
1081     loss_array.append(loss_testing)
1082
1083     accuracy_plot_array.append(accuracy_plot_training)
1084     accuracy_plot_array.append(accuracy_plot_validation)
1085     accuracy_plot_array.append(accuracy_plot_testing)
1086
1087     # printing error on training and testing dataset
1088     print('Error on validation dataset : ' + str(error_plot[min_error_index])
1089     + '%');
1090     print('Error on test dataset : ' + str(test_error*100/len(test_data)) +
1091     '%');
1092
1093     return weights, loss_array, accuracy_plot_array
1094
1095 #-----Main function-----
1096
1097 if __name__ == "__main__":
1098     numpy.random.seed(0)
1099     learning_rate = 0.0001
1100     N = 20000
1101     N_test = 2000
1102     lamda = 0.001          # regularization weightage parameter
1103     T = 2000
1104     iteration = 1000
1105     training_data, training_label, test_data, test_label, validation_data, validation_label
1106     = get_data(N, N_test)
1107
1108     weights, loss_array, accuracy_plot_array = fit(training_data, training_label,
1109     test_data, test_label, validation_data,
1110     validation_label)
1111
1112     plotlyGraphs(loss_array, ['Training Set', 'Validation Set', 'Test Set'], "Loss Function and
1113     Iterations")
1114     plotlyGraphs(accuracy_plot_array, ['Training Set', 'Validation Set', 'Test Set'], "Accuracy
1115     and Iterations")

```